



Deadlock handling

Prof J P Misra
BITS, Pilani




The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- Example1
 - ➔ System has 2 tape drives
 - ➔ P_1 and P_2 each hold one tape drive and each needs another one
- Example2
 - ➔ Semaphores A and B , initialized to 1

P_0
 $Wait(A);$
 $Wait(B);$

P_1
 $Wait(B)$
 $Wait(A)$

Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.



System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
 - Each resource type R_i has W_i instances
 - Each process utilizes a resource as follows:
 - 1.Request
 - 2.Use
 - 3.Release
- Request and release of resources is done through system calls



Deadlock Characterization

Deadlock can arise if **four** conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource(i.e. resource is non-sharable).
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .




Resource-allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- Request edge – directed edge $P_i \rightarrow R_j$
- Assignment edge – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances




- P_i requests instance of R_j




- P_i is holding an instance of R_j




Example of a Resource Allocation Graph



Resource Allocation Graph With A Deadlock



Resource Allocation Graph With A Cycle But No Deadlock





Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - ➔ If only one instance per resource type, then deadlock
 - ➔ If several instances per resource type, possibility of deadlock



Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
(Prevention or Avoidance)
- Allow the system to enter a deadlock state and then recover.



Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource(i.e. resource is non-sharable).
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .



Deadlock Prevention

- Ensure that one of the 4 necessary conditions for deadlock does not hold.
- Constraints the way requests can be made for resources
 - **Mutual exclusion**
 - not required for sharable resources;
 - Must hold for non-sharable resources.
 - ➔ Make non-sharable resources as sharable
 - ➔ some resources are intrinsically non sharable



Deadlock Prevention


- **Hold and wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - ➔ Require process to request and be allocated all its resources before it begins execution,
 - Need to know all the requirement in the beginning itself
 - ➔ Allow process to request resources only when the process has none.
 - If a process has resource and requires additional resources, it must release all resources it is holding and then make request
 - ➔ Low resource utilization; Starvation possible



Deadlock Prevention (Cont.)

- **No preemption – Approach - 1**
 - ➔ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - ➔ Preempted resources are added to the list of resources for which the process is waiting
 - ➔ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.
 - ➔ Cannot be generally applied to resources like printers, tape drives etc.

No preemption (cont-)





Deadlock Prevention (Cont.)

- **Circular wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration



Dead lock Handling approaches

-
- Prevention
 - Avoidance
 - Detection & Recovery



Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm
 - Dynamically finds the *resource-allocation state* in order to ensure that a circular-wait condition does not occur.
- Resource-allocation *state* is defined by
 - the number of available resources
 - the number of allocated resources,
 - the maximum requirement of the processes



Safe State

- When a process requests an available resource, the system needs to decide if immediate allocation leaves the system in a *safe state*.
- A state is safe if system can allocate resources to each process (Max requirement) in some order and still avoid deadlock
- System is in safe state if there exists a **safe sequence** of all processes.



-
- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe
 - if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on




Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state



Safe, Unsafe , Deadlock State



Example



	Maximum need	Allocation
P0	10	5
P1	4	2
P2	9	2


- Total no of resources = 12
- safe sequence < P1, P0, P2 >




Resource-Allocation Graph Algorithm

- *Claim edge* $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j ; represented by a dashed line.
- Claim edge converts to *request edge* when a process requests a resource.
- When a resource is released by a process, *assignment edge* reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.

Resource-Allocation Graph For Deadlock Avoidance



Unsafe State In Resource-Allocation Graph





Banker's Algorithm

- Multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- *Available*: Vector of length m . If $\text{Available}[j] = k$, there are k instances of resource type R_j available.
- *Max*: $n \times m$ matrix. If $\text{Max}[I, j] = k$, then process P_i may request at most k instances of resource type R_j .
- *Allocation*: $n \times m$ matrix. If $\text{Allocation}[I, j] = k$ then P_i is currently allocated k instances of R_j .
- *Need*: $n \times m$ matrix. If $\text{Need}[I, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$\text{Need}[I, j] = \text{Max}[I, j] - \text{Allocation}[I, j]$$



Safety Algorithm

-
1. Let $Work$ and $Finish$ be vectors of length m and n , respectively.
Initialize:
 1. $Work = Available$
 2. $Finish [i] = false$ for $i = 0, 1, \dots, n-1$.
 2. Find an index i such that both:
 1. (a) $Finish [i] == false$
 2. (b) $Need_i \leq Work$
 3. If no such i exists, go to step 4.
 3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.
 4. If $Finish [i] == true$ for all i , then the system is in a safe state.



Safe State example

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation Matrix

	R1	R2	R3
	9	3	6

Resource Vector

	R1	R2	R3
	0	1	1

Available Vector

(a) Initial state

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0



Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i .

If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \leq Need_i$ then go to step 2.
Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3.
Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:
 $Available = Available - Request_i$;
 $Allocation_i = Allocation_i + Request_i$;
 $Need_i = Need_i - Request_i$
 - If safe \Rightarrow the resources are allocated to P_i
 - If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored



Example of Banker's Algorithm

-
- 5 processes P_0 through P_4 ; 3 resource types A (10 instances), B (5 instances) and C (7 instances).
 - Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Example (Cont.)

- The content of the matrix Need is equal to Max – Allocation.

	<u>Need</u>				
	A	B	C	<u>Allocation</u>	<u>Max Available</u>
	A	B	C	A	B
P_0	7	4	3	0 1 0	7 5 3 3 3 2
P_1	1	2	2	2 0 0	3 2 2
P_2	6	0	0	3 0 2	9 0 2
P_3	0	1	1	2 1 1	2 2 2
P_4	4	3	1	0 0 2	4 3 3

- The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.

Example P_1 Request (1,0,2) (Cont.)

- Check that Request \leq Available that is, $(1,0,2) \leq (3,3,2) \Rightarrow true$.

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
P_0	0	1	0	7	4	3	2	3	0
P_1	3	0	2	0	2	0			
P_2	3	0	2	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?



Deadlock Detection


- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme



Single Instance of Each Resource Type

- Maintain *wait-for* graph
 - ➔ Nodes are processes
 - ➔ $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke an algorithm that searches for a cycle in the graph.

Resource-Allocation Graph and Wait-for Graph





Several Instances of a Resource Type

- *Available*: A vector of length m indicates the number of available resources of each type.
- *Allocation*: An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
- *Request*: An $n \times m$ matrix indicates the current request of each process. If $Request [i,j] = k$, then process P_i is requesting k more instances of resource type R_j .



Detection Algorithm

1. Let $Work$ and $Finish$ be vectors of length m and n , respectively.
Initialize:
 - (a) $Work = Available$
 - (b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then
 $Finish[i] = \text{false}$; otherwise, $Finish[i] = \text{true}$.
2. Find an index i such that both:
 - (a) $Finish[i] == \text{false}$
 - (b) $Request_i \leq Work$
If no such i exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = \text{true}$
go to step 2.
4. If $Finish[i] == \text{false}$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == \text{false}$, then P_i is deadlocked.

Example of Detection Algorithm

- Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time T_0 :

	<i>Allocation</i>			<i>Request</i>			<i>Available</i>		
	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
P_0	0	1	0	0	0	0	0	0	0
P_1	2	0	0	2	0	2			
P_2	3	0	3	0	0	0			
P_3	2	1	1	1	0	0			
P_4	0	0	2	0	0	2			

- Sequence $\langle P_0, P_2, P_3, P_4, P_1 \rangle$ will result in $Finish[i] = \text{true}$ for all i .



Example (Cont.)

- P_2 requests an additional instance of type C .

	<u>Request</u>		
	A	B	C
P_0	0	0	0
P_1	2	0	1
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- State of system?
 - Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes requests.
 - Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4 .



Detection-Algorithm Usage

- When, and how often, to invoke depends on:
 - ➔ How often a deadlock is likely to occur?
 - ➔ How many processes will need to be rolled back?
- Invoke the algorithm every time request for resource cannot be granted immediately.
- Once per hour
- When CPU utilization drops below 40%

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes → May be expensive
- Abort one process at a time until the deadlock cycle is eliminated → Invoke deadlock detection algorithm after every abort.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - Is process interactive or batch?



Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.



Thanks