

# Java Threads

# Outline

- **Creating a Java Thread**
- Synchronized Keyword
- Wait and Notify

# Basic threads

- We will create a thread that simply prints out a number 500 times in a row.

```
class MyThread extends Thread {  
    int i;  
    MyThread(int i) {  
        this.i = i;  
    }  
    public void run() {  
        for (int ctr=0; ctr < 500; ctr++) {  
            System.out.print(i);  
        }  
    }  
}
```

# Basic threads

- To show the difference between parallel and non-parallel execution, we have the following executable MyThreadDemo class

```
class MyThreadDemo {  
    public static void main(String args[]) {  
        MyThread t1 = new MyThread(1);  
        MyThread t2 = new MyThread(2);  
        MyThread t3 = new MyThread(3);  
        t1.run();  
        t2.run();  
        t3.run();  
        System.out.print("Main ends");  
    } }
```

# Basic threads

- Upon executing MyThreadDemo, we get output that looks like this

```
$ java MyThreadDemo  
1111111111...22222222.....33333.....Main ends
```


- As can be seen, our program first executed t1's run method, which prints 500 consecutive 1's.
- After that t2's run method, which prints consecutive 2's, and so on.
- Main ends appears at the last part of the output as it is the last part of the program.

# Basic threads

- What happened was serial execution, no multithreaded execution occurred
  - This is because we simply called MyThread's run() method

# Basic threads

## Program Flow



# Basic threads

- To start parallel execution, we call MyThread's start() method, which is built-in in all thread objects.

```
class MyThreadDemo {  
    public static void main(String args[]) {  
        MyThread t1 = new MyThread(1);  
        MyThread t2 = new MyThread(2);  
        MyThread t3 = new MyThread(3);  
        t1.start();  
        t2.start();  
        t3.start();  
        System.out.print("Main ends");  
    } }
```

# Basic threads


- When we run MyThreadDemo we can definitely see that three run methods are executing at the same time

```
> java MyThreadDemo
```

```
1111111111223311223211223333111Main ends11133322...
```

# Basic threads

## Program Flow



# Basic threads

- Note the appearance of the "Main ends" string in the middle of the output sequence

```
> java MyThreadDemo
```

```
1111111111223311223211223333111Main ends111333222...
```

- This indicates that the main method has already finished executing while thread 1, thread 2 and thread 3 are still running.

# Basic threads

- Running a main method creates a thread.
  - Normally, your program ends when the main thread ends.
- However, creating and then running a thread's start method creates a whole new thread that executes its run() method independent of the main method.

.....

# Pausing and Joining Threads

# Pausing threads

- Threads can be paused by the sleep() method.
- For example, to have MyThread pause for half a second before it prints the next number, we add the following lines of code to our for loop.

```
for (int ctr=0; ctr < 500; ctr++) {  
    System.out.print(i);  
    try {  
        Thread.sleep(500); // 500 milliseconds  
    } catch (InterruptedException e) { }  
}
```

- Thread.sleep() is a static method of class thread and can be invoked from any thread, including the main thread.

# Pausing threads

```
for (int ctr=0; ctr < 500; ctr++) {  
    System.out.print(i);  
    try {  
        Thread.sleep(500); // 500 milliseconds  
    } catch (InterruptedException e) {}  
}
```

- The InterruptedException is an unchecked exception that is thrown by code that stops a thread from running.
  - Unchecked exception causing lines must be enclosed in a try-catch, in this case, Thread.sleep().
- An InterruptedException is thrown when a thread is waiting, sleeping, or otherwise paused for a long time and another thread interrupts it using the interrupt() method in class Thread.

# Joins

- You can also make a thread stop until another thread finishes running by invoking a thread's join() method.
- For instance, to make our main thread to stop running until thread t1 is finished, we could simply say...

```
public static void main(String args[]) {  
    ...  
    t1.start();  
    try {  
        t1.join();  
    } catch (InterruptedException e) {}  
    t2.start();
```

- This would cause all the 1's to be printed out before thread 2 and thread 3 start.

.....

# Critical Section Problem

# The Critical Section Problem

- This section we will find out how to implement a solution to the critical section problem using Java
- Recall that only a single process can enter its critical section, all other processes would have to wait
  - No context switching occurs inside a critical section

# The Critical Section Problem

- Instead of printing a continuous stream of numbers, MyThread calls a print10() method in a class MyPrinter
  - print10() prints 10 continuous numbers on a single line before a newline.
- Our goal is to have these 10 continuous numbers printed without any context switches occurring.
  - Our output should be:

...  
1111111111  
1111111111  
2222222222  
3333333333  
1111111111  
...

# MyPrinter

- We define a class MyPrinter that will have our print10() method

```
class MyPrinter {  
    public void print10(int value) {  
        for (int i = 0; i < 10; i++) {  
            System.out.print(value);  
        }  
        System.out.println(""); // newline after 10 numbers  
    }  
}
```

# MyThread (revised)

- Instead of printing numbers directly, we use the print10() method in MyThread, as shown by the following

```
class MyThread extends Thread {  
    int i;  
    MyPrinter p;  
    MyThread(int i) {  
        this.i = i;  
        p = new MyPrinter();  
    }  
    public void run() {  
        for (int ctr=0; ctr < 500; ctr++) {  
            p.print10(i);  
        }  
    }  
}
```

# MyThreadDemo (revised)

- First, we will try to see the output of just a single thread running.

```
class MyThreadDemo {  
    public static void main(String args[]) {  
        MyThread t1 = new MyThread(1);  
        // MyThread t2 = new MyThread(2);  
        // MyThread t3 = new MyThread(3);  
        t1.start();  
        // t2.start();  
        // t3.start();  
        System.out.print("Main ends");  
    }  
}
```

*We comment out  
the other threads  
for now...*

# MyThreadDemo

- When we run our MyThreadDemo, we can see that the output really does match what we want.

```
> java MyThreadDemo  
1111111111  
1111111111  
1111111111  
1111111111  
1111111111  
1111111111  
...  
...
```

# MyThreadDemo (revised 2)

- However, let us try to see the output with other threads.

```
class MyThreadDemo {  
    public static void main(String args[]) {  
        MyThread t1 = new MyThread(1);  
        MyThread t2 = new MyThread(2);  
        MyThread t3 = new MyThread(3);  
        t1.start();  
        t2.start();  
        t3.start();  
        System.out.print("Main ends");  
    } }
```

*We run all three  
threads*

# MyThreadDemo

- Our output would look like the following

```
> java MyThreadDemo
1111111111
111112222222
1111
22233333332
...
...
```

- We do not achieve our goal of printing 10 consecutive numbers in a row
  - all three threads are executing the print10's method at the same time, thus having more than one number appearing on a single line.

# Critical Section Problem


- To achieve our goal there should not be any context switches happening inside print10()
  - Only a single thread should be allowed to execute inside print10()
- As can be seen, this is an example of the critical section problem
- To solve this, we can use some of the techniques discussed in our synchronization chapter

# Possible Solutions

- Busy Wait
- Wait and Notify
- Semaphores
- Monitors

# Synchronized keyword

- Now we will discuss Java's solution to the critical section problem
- Java uses a monitor construct
  - Only a single thread may run inside the monitor



*Even if P2 is calling method2(), it has to wait for P1 to finish*

# Synchronized keyword

- To turn an object into a monitor, simply put the synchronized keyword on your method definitions
  - Only a single thread can run any synchronized method in an object

```
class MyPrinter {  
    public synchronized void print10(int value) {  
        for (int i = 0; i < 10; i++) {  
            System.out.print(value);  
        }  
        System.out.println(""); // newline after 10  
    numbers  
    }  
}
```

# Still not working!

- However, even if we did turn our print10() method into a synchronized method, it will still not work

```
> java MyThreadDemo  
1111111111  
111112222222  
1111  
22233333332  
...
```

- To find out how to make it work, we need to first find out how the synchronized keyword works

# Intrinsic locks


- Every object in Java has an intrinsic lock



```
class MyPrinter {  
    synchronized print10() {  
    }}  
}
```


# Intrinsic locks

- When a thread tries to run a synchronized method, it first tries to get a lock on the object




# Intrinsic locks

- If a thread is successful, then it owns the lock and executes the synchronized code.




# Intrinsic locks

- Other threads cannot run the synchronized code because the lock is unavailable




# Intrinsic Locks

- Threads can only enter once the thread owning the lock leaves the synchronized method



# Intrinsic Locks

- Waiting threads would then compete on who gets to go next




# MyThreadDemo Failure

- So why doesn't our example work?
- Because each thread has its own copy of the MyPrinter object!

```
class MyThread extends Thread {  
    int i;  
    MyPrinter p;  
    MyThread(int i) {  
        this.i = i;  
        p = new MyPrinter(); // each MyThread creates  
        its own MyPrinter!  
    }  
    public void run() {  
        for (int ctr=0; ctr < 500; ctr++) {  
            p.print10(i);  
        }  
    }  
}
```

# Different Locks



# Important

- Recall our definition of synchronized methods
  - Only a single thread can run any synchronized method in an object
- Since each thread has its own MyPrinter object, then no locking occurs

# Solution

- So, to make it work, all threads must point to the same MyPrinter object

# Solution


- Note how all MyThreads now have the same MyPrinter object

```
class MyThread extends Thread {  
    int i;  
    MyPrinter p;  
    MyThread(int i, MyPrinter p) {  
        this.i = i; this.p = p  
    }  
    public synchronized void run() {  
        for (int ctr=0; ctr < 500;  
            ctr++) {  
            p.print10(i);  
        }  
    }  
}
```

```
class MyThreadDemo {  
    public static void main(String args[]) {  
        MyPrinter p = new  
        MyPrinter();  
        MyThread t1 = new  
        MyThread(1,p);  
        MyThread t2 = new  
        MyThread(2,p);  
        MyThread t3 = new  
        MyThread(3,p);  
        t1.start();  
        t2.start();  
        t3.start();  
    }  
}
```

# Solution

```
class MyThreadDemo {  
    main() {  
        MyPrinter p =  
            new MyPrinter();  
        MyThread t1 =  
            new MyThread(p);  
        MyThread t2 =  
            new MyThread(p);  
        MyThread t3 =  
            new MyThread(p);  
  
        "Main ends"  
    }  
}
```



# Locks and Doors

- A way to visualize it is that all Java objects can be doors
  - A thread tries to see if the door is open
  - Once the thread goes through the door, it locks it behind it,
  - No other threads can enter the door because our first thread has locked it from the inside
  - Other threads can enter if the thread inside unlocks the door and goes out
  - Lining up will only occur if everyone only has a single door to the critical region

# Synchronized methods

- Running MyThreadDemo now correctly shows synchronized methods at work

```
>java MyThreadDemo  
1111111111  
1111111111  
1111111111  
2222222222  
3333333333  
...
```

# Remember

- Only a single thread can run any synchronized method in an object
  - If an object has a synchronized method and a regular method, only one thread can run the synchronized method while multiple threads can run the regular method
- Threads that you want to have synchronized must share the same monitor object

# Synchronized blocks

- Aside from synchronized methods, Java also allows for synchronized blocks.
- You must specify what object the intrinsic lock comes from

# Synchronized blocks

- Our print10() implementation, done using synchronized statements, would look like this

```
class MyPrinter {  
    public void print10(int value) {  
        synchronized(this) {  
            for (int i = 0; i < 10; i++) {  
                System.out.print(value);  
            }  
            System.out.println(""); // newline after  
10 numbers  
        }  
    }  
}
```

*We use a MyPrinter object for the lock, replicating what the synchronized method does*

# Synchronized blocks

- Synchronized blocks allow for flexibility, in the case if we want our intrinsic lock to come from an object other than the current object.
- For example, we could define MyThread's run method to perform the lock on the MyPrinter object before calling print10

```
class MyThread extends Thread {  
    MyPrinter p;  
    ...  
    public void run() {  
        for (int i = 0; i < 100; i++) {  
            synchronized(p) {  
                p.print10(value);  
            }  
        }  
    }  
}
```

*Note that any lines before and after our synchronized block can be executed concurrently by threads*

# Multiple locks

- Also, the use of the synchronized block allows for more flexibility by allowing different parts of code to be locked with different objects.


# Multiple locks

- For example, consider a modified MyPrinter which has two methods, which we'll allow to run concurrently

```
class MyPrinter {  
    Object lock1 = new Object();  
    Object lock2 = new Object();  
    public void print10(int value) {  
        synchronized(lock1) {  
            for (int i = 0; i < 10; i++) {  
                System.out.print(value);  
            }  
            System.out.println(""); // newline after 10 numbers  
        }  
    }  
    public int squareMe(int i) {  
        synchronized (lock2) {  
            return i * i;  
        }  
    }  
}
```

The blocks inside print10() and squareMe() can run at the same time because they use different objects for their locking mechanism

Synchronization still applies. For example only a single thread can run the synchronized block in squareMe() at any point in time



Any object in Java can be used as a lock

Thread 5 and 8 are allowed at the same time in synchblock as they use different locks

Only one thread in synchronized block. As many threads everywhere else

# Outline

- Creating a Java Thread
- Synchronized Keyword
- **Wait and Notify**
- High Level Cuncurrency Objects

# Synchronization/`wait` and `notify`

- `wait` is a method of the `Object` class. It causes the calling method to wait (blocked) until notified (when another thread calls `notify` or `notifyAll`)
- While waiting, a thread **relinquishes** its object locks. This gives other thread a chance to access the object.
- `notify` or `notifyAll` are all methods of the `Object` class.
  - `notify` randomly selects a thread among those waiting for an object lock (inside the object) and unblocks it.
  - `notifyAll` unblocks all threads waiting for an object lock (inside the object). Preferred because it reduces the probability of deadlock (Exercise: Why is this?).
- Note: `wait`, `notify` and `notifyAll` should only be called by a thread that is the owner of this object's monitor (within synchronized method).

# Synchronization/**wait** and **notify**

- In our example, a thread calls **notifyAll** when it is done with a `Object`.

```
public synchronized void transfer(int from, int to,  
                                 int amount)  
{   while (accounts[from] < amount)  
   {     try {   wait(); }  
       catch(InterruptedException e) {}  
   }  
   ...  
   notifyAll();  
...}
```

SynchBankTest.java

# Synchronization/Deadlock

Deadlock occurs when a number of threads waiting for each other

Example:

Account 1: \$2,000

Account 2: \$3,000

Thread 1: Transfer \$3,000 from account 1 to account 2

Thread 2: Transfer \$4,000 from account 2 to account 1

It is the responsibility of programmer to avoid deadlocks.

# Summary

- A thread is a path of execution that executes within a process. When you run a Java program, the main() method runs in a thread. The main() method can start other threads.
- A process terminates when all its non-daemon threads run to completion.
- A thread is created by writing a class that implements java.lang. Runnable and associating an instance of this class with a new java.lang.Thread object. The new Thread is initially in the born state.
- Invoking the start() method on a Thread object starts the thread and places it in its appropriate runnable queue, based on its priority.

# Summary

- Java uses a preemptive scheduling mechanism where threads of higher priority preempt running threads of a lower priority. However, the actual behavior of threads also relies on the underlying platform.
- A thread can also be written by writing a class that extends `java.util.TimerTask` and associating an instance of this class with a `java.util.Timer` object. This type of thread is useful when performing scheduled tasks.

# Summary

- The synchronized keyword is used to synchronize a thread on a particular object. When a thread is synchronized on an object, the thread owns the lock of the object. Any other thread attempting to synchronize
- Care needs to be taken when synchronizing threads, since deadlock may occur. Ordering locks is a common technique for avoiding deadlock.
- The wait() and notify() methods from the Object class are useful when multiple threads need to access the same data in any type of producer/consumer scenario.